

This is part of a three part hand-in to the M1-exam.

Link to this colab is <https://colab.research.google.com/drive/1i7VYq8SgB5dFA8jNcMaVJBBHqsD43exl>

Preprocessing

Importing libraries and fetching data

```
## Get data
!wget -c https://github.com/bande15/SDS_M1_EXAM/raw/master/combine_data.csv -O df.csv

## Mute Warnings
from warnings import simplefilter
simplefilter(action='ignore', category=(FutureWarning, DeprecationWarning, Warning))

#Import General Libraries
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import xgboost as xgb
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier

#Install and import special libraries
!pip install catboost
from catboost import CatBoostClassifier
```

Reading and processing data; The analysis is trying to identify and predict position from physical stats, so all other before missing values are dropped. Data is then converted to metric.

```
df = pd.read_csv('df.csv')
df_pre = df.drop(['Player', 'Pfr_ID', 'Year', 'AV'], 1)
df_pre = df_pre.dropna(how='any', axis=0)
df = df.drop(['Player', 'Pfr_ID', 'Year', 'AV', 'Round', 'Pick', 'Team'], 1)
a = len(df)
df = df.dropna(how='any', axis=0)
b = a - len(df)
a = len(df)
print(b, 'observations dropped due to missing values.', a, 'values left.')
del a, b
```

☞ 3333 observations dropped due to missing values. 2885 values left.

We will now count occurrences for each unique position before and after observations with missing data is dropped.

```
df_old = pd.read_csv('df.csv')
a = df_old.Pos.value_counts()
b = df.Pos.value_counts()
print(a, b)
```

```

↳ WR      857
   CB      630
   RB      540
   DE      487
   DT      463
   OT      460
   OLB     424
   OG      365
   QB      350
   TE      337
   ILB     276
   FS      229
   SS      213
   C       171
   P       120
   FB      117
   K        85
   S        27
   EDGE     23
   LS       20
   G        14
   NT        3
   LB        3
   DB        2
   OL        2
Name: Pos, dtype: int64 CB      311
DE      279
WR      275
OT      273
DT      253
RB      245
OLB     240
OG      224
TE      194
ILB     130
FS      123
C       115
SS      107
FB       77
QB       12
G         8
EDGE      8
S         7
LS         2
LB         1
OL         1
Name: Pos, dtype: int64

```

From the above lists of observations per position it becomes clear that some positions have very few or no cohe only does this prevent a 'coherent' analysis, but the option of simply scrappinf the observations, or konverting mi yield truthful results. This is a essential fundamentiil problem with the following analysis, and it will be discussed handout. The short version is, excludidng positions will not only prevent the models from predicting those positio associate some combination of stats with other positions, this yielding false results.

```

# Metric Conversion
def to_cm(F):
    return F*2.54
def to_kg(F):
    return F*0.45359237

```

```
df[['Ht', 'Vertical', 'BroadJump']] = df[['Ht', 'Vertical', 'BroadJump']].apply(to_cm)
df[['Wt']] = df[['Wt']].apply(to_kg)
df_pre[['Ht', 'Vertical', 'BroadJump']] = df_pre[['Ht', 'Vertical', 'BroadJump']].apply(to_cm)
df_pre[['Wt']] = df_pre[['Wt']].apply(to_kg)
df.info()
df.head()
```



With the newly cleaned and metrified dataframe, we can now try to make some explorative data analysis.

First off, we can look at the different teams, asking the question:

How different are the teams?

This can be examined by grouping for the 'Team' variable, and taking the mean of relevant variables.

```
df_pre.groupby('Team')['Ht', 'Wt', 'Forty', 'Vertical', 'BenchReps', 'BroadJump', 'Cone', 'Shuttle', 'Rou
```



From above output it looks like the different teams *generally* have similar statistics. This might imply that the team has similar physical stats in their player picks. Thus it also becomes interesting to look at the different positions.

Similar to the above code, we can group by the position variable, 'Pos', and take the mean of the remaining relevant variables.

```
df_pre.groupby('Pos')['Ht', 'Wt', 'Forty', 'Vertical', 'BenchReps', 'BroadJump', 'Cone', 'Shuttle', 'RoundTrip']
```



From above output it becomes very clear, that players playing different positions also have different physical traits. An extreme example is the positions CB and OG. Looking at those two specific positions, it becomes clear that CB is lighter, can't take as many bench reps, but they can jump higher.

Those results point towards the following hypothesis:

Player positions implies different skills, thus requiring different physical attributes.

▼ Unsupervised learning

Focusing on above hypothesis, we can try to use unsupervised learning to see, if the data itself is capable of displaying players' positions based on their physical stats.

Thus our first step will be to make a new dataframe, `df_u`, based on the original dataframe, `df`, as defined earlier by transforming the 'Pos' variable into dummy variables using the `get_dummies` function, specifying the `drop_first`:

```
df_u = pd.get_dummies(df, drop_first=True)
df_u.head()
```



To avoid overspecification we reduce the dataset to 'fundamentals' by excluding stats that could be derived. That will be picked first for a specific position because of physical stats, so including both physical stats and position, may lead to overspecification. Before running unsupervised machine learning, we reduce the data to fundamental physical stats.

We will now be scaling our variables, by using the StandardScaler function. Scaling is important when using dist KMeans is. It can also be a good idea to scale your data, if the different variables have large differences in their After scaling we use the elbow method on the inertias plot, to determine a rightful amount of clusters.

```
ks = range(1, 15)
inertias = []
scale = StandardScaler()
sample = df_u
scale.fit(sample)
samples = scale.transform(sample)

for i in ks:
    model = KMeans(n_clusters=i)
    model.fit(samples)
    inertias.append(model.inertia_)

plt.figure(figsize=(20,12))
plt.plot(ks, inertias, '-o')
plt.xlabel('number of clusters, i')
plt.ylabel('inertia')
plt.xticks(ks)
plt.show()
```



Looking at above plot we determine k clusters = 2, thus we want to use a total of 2 clusters.

```
model = KMeans(n_clusters=2)
model.fit(samples)
labels = model.predict(samples)
```

Now we want to use the PCA algorithm on our data. Principal components analysis is a helpful technique in order dimension reduction. Thus it helps us by showing how many PCA features we will include in further analysis.

```
#Fitting the PCA algorithm with our Data
pca = PCA().fit(samples)
#Plotting the Cumulative Summation of the Explained Variance
plt.figure()
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('Number of Components')
plt.ylabel('Variance (%)') #for each component
plt.title('PCA component selection')
plt.show()
```



Looking at above plot of the explained variance, we choose 17 principal components. In comparison there are 2 a few of them being practically identical.

Using 17 principal components, the clusters can now be made and plotted.

```
pca = PCA(n_components=17)
sample = pca.fit_transform(samples)
xs = sample[:,0]
ys = sample[:,1]
plt.scatter(xs, ys, c=labels)
plt.axis('equal')
plt.show()
```



Looking at above plot it seems like the algorithm is capable of splitting the observations into two well-defined and
Following this, it could be very interesting to see, how the different player positions are spread throughout the cl
This can be examined using the crosstab function. In the following we have set normalize='index'. What this doe
the different rows, thus summarizing the two clusters for any given position equals 1.

```
lab = pd.DataFrame(labels)
lab.columns = ['Clusters']
df_up = df_u.join(lab)
pd.crosstab(df.Pos, df_up.Clusters, normalize='index')
```



Even though most roles are represented more in one cluster than the other, even at a ratio of 2:1, it doesn't look particularly good at defining the different positions, thus the clusters are not particularly informative.

▼ Supervised learning

From unsupervised learning, it did not appear like the clusters were capable of splitting the players into different position. In the following code, we will try to use different supervised ML algorithms to see if this works any better position a player has, based on how their physical traits are.

Thus the question becomes: can we instead predict position by using some supervised machine learning model?

To run a supervised data analysis, the data is first converted to numeric, then scaled, and split into test and train variables separated from the datasets.

```
supervised = df
supervised.head()
```



▼ Changing the position variable into a numerical variable

We want to use the position variable (Pos) as our target variable. Because the variable is categorical it is first converted to a numerical variable.

```
posdict = {}
a = supervised.Pos.unique()
b = 0
while b < len(a):
    posdict[a[b]] = b
    b += 1
supervised['PosN'] = supervised['Pos'].map(posdict)
```

▼ Splitting data into test- and train-sets

We now want to split the X- and y-variables into a train- and test-set. We do this by using the train_test_split function from sklearn.model_selection. Before doing this, we want to use a scaler (the StandardScaler() function).

Afterwards we split the data into a training and a test set. For this purpose we use a test size of 25%, and we set a random state value (42), so we can easily replicate our results.

```
encoder = LabelEncoder()
scaler = StandardScaler()
dataset = scaler.fit_transform(supervised.drop(['Pos', 'PosN'], 1))
target = encoder.fit_transform(supervised['PosN'])
train_data, test_data, train_target, test_target = train_test_split(dataset, target, test_size=0.25, random_state=42)
```

Now that the data has been split into a training and a test set, and it is scaled, we can use different models and see how the models perform.

▼ Linear Regression

```
# K-fold cross-validation
model1 = LinearRegression()
scores1 = cross_val_score(model1, train_data, train_target, cv = 5)
print("An average cross validation score of {}".format(np.mean(scores1)))

# Model training
model1.fit(train_data, train_target)

# Model performance
print('Model performance:', model1.score(test_data, test_target))
```



Unsurprisingly the model has a very poor performance, and it will not be used any further.

▼ Logistic Regression

```
# K-fold cross-validation
model2 = LogisticRegression()
scores2 = cross_val_score(model2, train_data, train_target, cv = 5)
print("An average cross validation score of {}".format(np.mean(scores2)))

# Model training
model2.fit(train_data, train_target)

# Model performance
print('Model performance:', model2.score(test_data, test_target))
```



LogisticRegression is much better at predicting positions, but it still only gets roughly half correct.

▼ Decision Tree model

```
model3 = DecisionTreeClassifier()
scores3 = cross_val_score(model3, train_data, train_target, cv = 5)
print("An average cross validation score of {}".format(np.mean(scores3)))

# Model training
model3.fit(train_data, train_target)

# Model performance
print('Model performance:', model3.score(test_data, test_target))
```



DecisionTree is slightly worse than LogisticRegression, but still much better than LinearRegression.

▼ CatBoost

```
model4 = CatBoostClassifier(logging_level='Silent')
scores4 = cross_val_score(model4, train_data, train_target, cv = 5)
```

```
print("An average cross validation score of {}".format(np.mean(scores4)))

# Model training
model4.fit(train_data, train_target)

# Model performance
print('Model performance:', model4.score(test_data, test_target))
```



CatBoostClassifier is so far the best model with a performance of roughly 0.57.

▼ XGBoost

```
model5 = xgb.XGBClassifier()
scores5 = cross_val_score(model5, train_data, train_target, cv = 5)
print("An average cross validation score of {}".format(np.mean(scores5)))

# Model training
model5.fit(train_data, train_target)

# Model performance
print('Model performance:', model5.score(test_data, test_target))
```



The XGBoostClassifier appears to be slightly worse compared to CatBoostClassifier. It performs better than Log DecisionTree though.

▼ Predicting the models

We call the .predict function on test_data in the models (except LinearRegression, as it was absolutely garbage) predictions in y_pred2, y_pred3, y_pred4 and y_pred5.

```
y_pred1 = model1.predict(test_data)
y_pred2 = model2.predict(test_data)
y_pred3 = model3.predict(test_data)
y_pred4 = model4.predict(test_data)
y_pred5 = model5.predict(test_data)
```

▼ Evaluating the models using classification_report and confusion_matrix

Using the confusion matrixes, correctly predicted positions can be seen on the diagonal. Thus a perfect model would have 1 on the diagonal and nowhere else. Looking at the CatBoostClassifier it becomes clear, that the model is the best of the models, at predicting which positions a player is drafted to. In certain positions, however, it looks like the model is predicting them.

Overall it looks like the CatBoost model is good at predicting which position a player has - and in situations where it predicts right, it generally looks like the model predicts 2-3 different positions. This behaviour might be explained, requiring almost identical physical traits, which also makes it hard for the model to distinguish the positions from one another. In consideration, it does appear that the players' physical traits are deterministic to which positions the players

▼ Logistic Regression

```
print(classification_report(test_target, y_pred2))  
print(confusion_matrix(test_target, y_pred2))
```



▼ Decision Tree

```
print(classification_report(test_target, y_pred3))  
print(confusion_matrix(test_target, y_pred3))
```



▼ **CatBoostClassifier**

```
print(classification_report(test_target,y_pred4))  
print(confusion_matrix(test_target,y_pred4))
```



▼ **XGBoostClassifier**

```
print(classification_report(test_target,y_pred5))  
print(confusion_matrix(test_target,y_pred5))
```



Discussion

The critical flaw in this analysis is the missing data. We cannot create a coherent analysis of only some positions are not independent, but defined by their relation to other positions. Therefore trying to define the positions based without the other positions are ultimately futile. Consider the following example: If weight was very heavily correlated with LS (Long snapper) and DT (Defensive tackle), then a high 'Wt' value might indicate **either** DT or LS, but if the LS has a low number of observations, the calculated correlation with weight and the DT positions will be unnatural. The number of discussion twists the entire analysis, and including positions with low number of observations undermines the analysis. In other words, a coherent analysis is not possible from the given dataset.

If, for the sake of the exercise, the data is treated as valid, we can see from the confusion matrices that roughly 80% of the predictions are correct, and the rest of the 'mispredicted' are not randomly distributed, but clustered in positions similar to the correct ones.

Conclusion

The models have not achieved a theoretical coherency, but have nevertheless managed to predict the correct position more often than chance.

